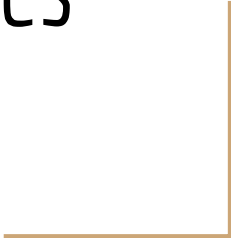


Graph Analytics at Scale using SparkSQL and GraphFrames



Overview

- What is Spark?
- SparkSQL & GraphFrames
- What is meant by “scale”?
- Graph Algorithms in GraphFrames
- Walkthrough: the Alternating Algorithm (Connected Components)
- Demo

What is Spark?

From the Apache Spark website:

“Unified analytics engine for large-scale data processing”

For our purposes a few things you need to know:

- Defines an API (in several programming languages) for manipulating data on distributed computing systems using *drivers* and *executors*
- Core data abstraction in Spark is a resilient-distributed dataset (RDD)

Spark SQL

Spark SQL sits on top of Spark core to provide a higher-level data abstraction called a DataFrame.

- Provides data users a familiar interface to query data (SQL) in addition to the Spark API
- Structure implied by DataFrames allows for additional optimizations and these optimizations are used regardless of which programming language you use

GraphFrames

Graphframes is an open-source project (not part of Spark Core) which provides DataFrame-based graphs.

- Utilizes DataFrames instead of RDDs to execute a subset of graph algorithms
- Implements a core set of graph algorithms on data at scale
- Functionality is similar to GraphX
 - “GraphX is to RDDs as GraphFrames are to DataFrames.”

What is meant by “scale”?

1. From a hardware perspective, you can improve the performance of your application by:
 - Using more powerful machines (vertical scaling)
 - Using more machines (horizontal scaling)
2. From a software perspective, the algorithms you choose to implement can execute on graphs with billions of nodes and edges

So... does GraphFrames “scale”?

1. From a hardware perspective, you can improve the performance of your application by:
 - Using more powerful machines (vertical scaling)
 - Using more machines (horizontal scaling)
2. From a software perspective, the algorithms you choose to implement can execute on graphs with billions of nodes and edges

So... does GraphFrames “scale”?

1. From a hardware perspective, you can improve the performance of your application by:
 - Using more powerful machines (vertical scaling)
 - Using more machines (horizontal scaling)
2. From a software perspective, the algorithms you choose to implement can execute on graphs with billions of nodes and edges



So... does GraphFrames “scale”?

1. From a hardware perspective, you can improve the performance of your application by:
 - Using more powerful machines (vertical scaling)
 - Using more machines (horizontal scaling)
2. From a software perspective, the algorithms you choose to implement can execute on graphs with billions of nodes and edges



GraphFrames

The core data structure in the GraphFrames library is, unsurprisingly, a GraphFrame. It's a class that is instantiated with two SparkSQL DataFrames:

- Vertices (Nodes) - A DataFrame of all the unique vertices/nodes of your graph
- Edges (Linkages) - A DataFrame of all the edges/linkages between the vertices/nodes in your graph

GraphFrames

```
# Vertex DataFrame
```

```
v = spark.createDataFrame([  
  ("a", "Alice", 34),  
  ("b", "Bob", 36),  
  ("c", "Charlie", 30),  
  ("d", "David", 29),  
  ("e", "Esther", 32),  
  ("f", "Fanny", 36),  
  ("g", "Gabby", 60)  
], ["id", "name", "age"])
```

```
# Edge DataFrame
```

```
e = spark.createDataFrame([  
  ("a", "b", "friend"),  
  ("b", "c", "follow"),  
  ("c", "b", "follow"),  
  ("f", "c", "follow"),  
  ("e", "f", "follow"),  
  ("e", "d", "friend"),  
  ("d", "a", "friend"),  
  ("a", "e", "friend")  
], ["src", "dst", "relationship"])
```

```
# Create a GraphFrame
```

```
g = GraphFrame(v, e)
```

GraphFrames

```
# Vertex DataFrame
```

```
v = spark.createDataFrame([  
  ("a", "Alice", 34),  
  ("b", "Bob", 36),  
  ("c", "Charlie", 30),  
  ("d", "David", 29),  
  ("e", "Esther", 32),  
  ("f", "Fanny", 36),  
  ("g", "Gabby", 60)  
], ["id", "name", "age"])
```

```
# Edge DataFrame
```

```
e = spark.createDataFrame([  
  ("a", "b", "friend"),  
  ("b", "c", "follow"),  
  ("c", "b", "follow"),  
  ("f", "c", "follow"),  
  ("e", "f", "follow"),  
  ("e", "d", "friend"),  
  ("d", "a", "friend"),  
  ("a", "e", "friend")  
], ["src", "dst", "relationship"])
```

```
# Create a GraphFrame
```

```
g = GraphFrame(v, e)
```

Graph Algorithms in GraphFrames

- Breadth-first Search
- Connected Components
 - Strongly Connected Components
- Label Propagation Algorithm
- PageRank
- Shortest Path
- Triangle Count

See the [docs](#) for more details on each of these.

Custom Graph Algorithms

GraphFrames provides primitives for developing graph algorithms

- aggregateMessages API
- Pregel API

Example: [Belief propagation](#)

Connected Components

For a given node, what other nodes are connected to it, either directly or indirectly?

Connected Components algorithms make explicit connections between nodes in a graph by assigning each node a cluster ID. Nodes with the same cluster IDs are connected, directly or indirectly.

Connected Components - Motivation


Nodes <i>ID</i>	Edges	
	<i>From</i>	<i>To</i>
1	1	8
2	5	8
3	8	9
4	7	8
5	3	7
6	2	3
7	3	6
8	4	6
9		

Connected Components - Motivation



Nodes		Edges	
ID		From	To
1		1	8
2		5	8
3		8	9
4		7	8
5		3	7
6		2	3
7		3	6
8		4	6
9			


Connected Components - Motivation



Nodes ID	Edges	
	From	To
1	1	8
2	5	8
3	8	9
4	7	8
5	3	7
6	2	3
7	3	6
8	4	6
9		

Connected Components - Motivation

From the edges table, how would we know that nodes 1 and 4 are connected?




Nodes ID	Edges	
	From	To
1	1	8
2	5	8
3	8	9
4	7	8
5	3	7
6	2	3
7	3	6
8	4	6
9		

Connected Components - Motivation

We know if nodes are connected only if there is a direct connection in the edge table. Any indirect connections that may exist can be determined via a connected components algorithm.


One approach is to reformulate the edge table until all nodes that are connected, either directly or indirectly, are made to be explicitly connected directly to a common node.

Connected Components - Motivation



Nodes ID	Edges	
	From	To
1	1	8
2	5	8
3	8	9
4	7	8
5	3	7
6	2	3
7	3	6
8	4	6
9		


Connected Components - Motivation



Nodes	Edges	
	From	To
1	1	2
2	1	3
3	1	4
4	1	5
5	1	6
6	1	7
7	1	8
8	1	9
9		

Connected Components - Motivation

Cluster ID: 1



Nodes ID	Edges	
	From	To
1	1	2
2	1	3
3	1	4
4	1	5
5	1	6
6	1	7
7	1	8
8	1	9
9		

Connected Components


The Connected Component algorithm implemented in GraphFrames is called the Alternating Algorithm, as defined in the Google research paper from 2014:

[Connected Components in MapReduce and Beyond](#)

This algorithm works by iteratively reformulating the edge table by alternating two operations - the Large Star operation and the Small Star operation - until convergence.


Alternating Algorithm - Initialization

Randomly assign integer IDs to all nodes in graph



Alternating Algorithm - Initialization

Randomly assign integer IDs to all nodes in graph




Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)


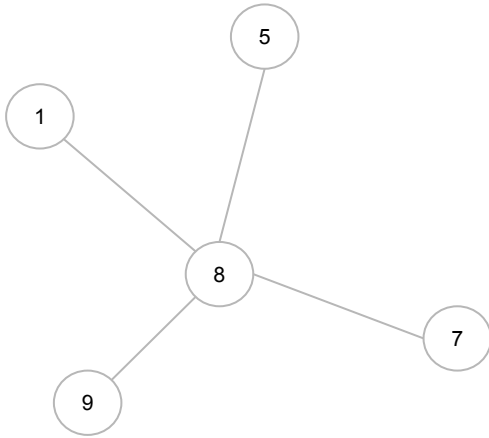
Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)



Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)





Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)





Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)





Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)





Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)





Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)





Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)





Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)





Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)



Large Star Operation

For each node in the graph, connect all strictly **larger** neighbors to the **min** neighbor (including **self**)




Large Star Operation

The large star operation has two theoretical guarantees:

1. Preserves connectivity of components
2. Never increases the number of edges in the graph



Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor





Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor





Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor





Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor





Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor





Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor





Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor





Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor





Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor



Small Star Operation

For each node in the graph, connect all **smaller** neighbors (and **self**) to the **min** neighbor




Small Star Operation

The small star operation has the same theoretical guarantees as the large star operation:


1. Preserves connectivity of components
2. Never increases the number of edges in the graph

Alternating Algorithm - Example



Nodes		Edges	
ID		From	To
1		1	8
2		5	8
3		8	9
4		7	8
5		3	7
6		2	3
7		3	6
8		4	6
9			


Alternating Algorithm - Example



Nodes		Edges	
ID		From	To
1		1	8
2		5	8
3		8	9
4		7	8
5		3	7
6		2	3
7		3	6
8		4	6
9			

Alternating Algorithm - Example



Step 1 - Large Star (1)



Nodes	Edges	
	From	To
1	1	9
2	1	8
3	5	8
4	3	8
5	2	7
6	2	3
7	2	6
8	4	6
9		

Alternating Algorithm - Example


Step 2 - Small Star (1)



Nodes	Edges	
	From	To
1	1	9
2	1	8
3	1	5
4	1	3
5	2	7
6	2	3
7	2	6
8	2	4
9		

Alternating Algorithm - Example


Step 3 - Large Star (2)



Nodes	Edges	
	From	To
1	1	9
2	1	8
3	1	5
4	1	3
5	2	7
6	2	3
7	2	6
8	2	4
9		

Alternating Algorithm - Example


Step 4 - Small Star (2)



Nodes	Edges	
	From	To
1	1	9
2	1	8
3	1	5
4	1	3
5	2	7
6	2	5
7	2	6
8	2	4
9		

Alternating Algorithm - Example

Step 5 - Large Star (2)



Nodes	Edges	
	From	To
1	1	9
2	1	8
3	1	5
4	1	3
5	1	7
6	1	5
7	1	6
8	1	4
9		

GraphFrames - Large Star and Small Star

```
...while (!converged) {  
  ...// Large-star step  
  ...// compute min neighbors (including self-min)  
  ...val minNbrs1 = minNbrs(ee) // src >= min_nbr  
  ...persist(intermediateStorageLevel)  
  ...// connect all strictly larger neighbors to the min neighbor (including self)  
  ...ee = skewedJoin(ee, minNbrs1, broadcastThreshold, logPrefix)  
  ...select(col(DST).as(SRC), col(MIN_NBR).as(DST)) // src > dst  
  ...distinct()  
  ...persist(intermediateStorageLevel)  
  
  ...// small-star step  
  ...// compute min neighbors (excluding self-min)  
  ...val minNbrs2 = ee.groupBy(col(SRC)).agg(min(col(DST)).as(MIN_NBR), count("*").as(CNT)) // src > min_nbr  
  ...persist(intermediateStorageLevel)  
  ...// connect all smaller neighbors to the min neighbor  
  ...ee = skewedJoin(ee, minNbrs2, broadcastThreshold, logPrefix)  
  ...select(col(MIN_NBR).as(SRC), col(DST)) // src <= dst  
  ...filter(col(SRC) != col(DST)) // src < dst  
  ...// connect self to the min neighbor  
  ...ee = ee.union(minNbrs2.select(col(MIN_NBR).as(SRC), col(SRC).as(DST))) // src < dst  
  ...distinct()  
}
```

GraphFrames - Convergence

The Alternating Algorithm converges when the sum of the SRC column values does not change from one iteration (Large Star step and Small Star step) to the next.

GraphFrames - Convergence

```
...//test convergence

...//Taking the sum in DecimalType to preserve precision.
...//We use 20 digits for long values and Spark SQL will add 10 digits for the sum.
...//It should be able to handle 200 billion edges without overflow.
...val (currSum, cnt) = ee.select(sum(col(SRC).cast(DecimalType(20, 0))), count("*")).rdd
...  .map { r =>
...    (r.getAs[BigDecimal](0), r.getLong(1))
...  }.first()
...  if (cnt != 0L && currSum == null) {
...    throw new ArithmeticException(
...      s"""
...        |The total sum of edge src IDs is used to determine convergence during iterations.
...        |However, the total sum at iteration $iteration exceeded 30 digits (1e30),
...        |which should happen only if the graph contains more than 200 billion edges.
...        |If not, please file a bug report at https://github.com/graphframes/graphframes/issues.
...        |"""
...      .stripMargin)
...  }
...  logInfo(s"$logPrefix Sum of assigned components in iteration $iteration: $currSum.")
...  if (currSum == prevSum) {
...    //This also covers the case when cnt = 0 and currSum is null, which means no edges.
...    converged = true
...  } else {
...    prevSum = currSum
...  }

...  iteration += 1
...}
```

Demo